

ML2 Project

Eric Jonas, Caroline Graebel, Naim Iskandar Zahari, Kseniia Holovchenko

2025-01-06

Tip: For the optimal experience view the report in HTML format.

Introduction

Understanding rent prices in Berlin is essential for anyone looking to rent an apartment in the city, given its dynamic housing market and the growing demand for affordable living spaces. To address this, our project focuses on analyzing historical rental data from Berlin in 2020, aiming to uncover insights and develop predictive models that estimate rental prices based on key property attributes and evaluate them.

The following dataset, available at Kaggle, contains information about rental properties listed on ImmoScout24. The data was scraped from ImmoScout24 between February and October 2020 and includes listings from all German federal states. For our project, we will focus exclusively on Berlin and train a model to predict either the `totalRent` or the `basePrice`.

Data Description

The data provided in the Kaggle dataset consists of approximately 268,000 rental listings in Germany. In the dataset, the information from the rental listings are structured in a CSV file with 49 columns. The columns include:

- Geographical information (state, city name and zip code)
- Property information (construction year, flat type, heating type and number of rooms)
- Total rent, base rent and service charge (in Euros)

Objectives

- **Data Understanding and Cleaning:**

Before training a model, it is essential to thoroughly understand and clean the data. To achieve this, we will provide a concise visual overview of the dataset to highlight its key characteristics and identify potential issues.

- **Predict Missing Values:**

Because many columns in our data contain missing values, one of our goals is to impute these values and use the resulting data instead of the raw data.

- **Model Selection and Comparison:**

We aim to train a series of tree-based models on the dataset to predict the `baseRent` of an apartment, starting from simple decision trees that can provide explainable predictions and gradually progressing to more complex models such as Random Forests. Additionally, we will compare the performance of these tree-based methods with Support Vector Machine (SVM) regression to determine which approach yields the best results.

Data Overview and Cleaning

Historical rental listing data for Berlin is crucial for achieving the objectives of this project. The ImmoScout24 dataset provides rental listings from all of Germany. To begin, we need to extract data specific to Berlin from the full dataset. This extracted data contains numerous missing values and redundant entries, making it essential to first thoroughly understand and clean the dataset before proceeding.

Data Preparation

After filtering for Berlin listings, more than 10,000 entries remain. The next step is to remove irrelevant fields from the dataset, which will help with the development of machine learning models in subsequent stages. Identifying and removing these irrelevant fields requires a thorough understanding of the dataset. A brief explanation of each column can be found on Kaggle. Below is a list of the fields removed, along with the reasons for their exclusion:

Table 1: Columns and Reasons for Removal

Column	Reason for Removal
regio1	Redundant, as it contains only 'Berlin.'
geo_bln	Redundant, as it contains only 'Berlin.'
geo_krs	Redundant, as it contains only 'Berlin.'
regio2	Redundant, as it contains only 'Berlin.'
street	Redundant information; contains the same data as street_plain , but of lower quality.
telekomTvOffer	Irrelevant, as all listings can offer this service.
telekomHybridUploadSpeed	Unnecessary, with mostly missing (NA) values.
scoutId	Unique ID, not useful for analysis.
yearConstructedRange	Artificial column, not relevant to the analysis.
streetPlain	Would require significant cleaning or conversion into longitude and latitude.
houseNumber	Requires a corresponding street to provide meaningful context.
baseRentRange	Artificial column, not needed for analysis.
livingSpaceRange	Artificial column, not needed for analysis.
noRoomsRange	Artificial column, not needed for analysis.
electricityKwhPrice	Deprecated since 2020; irrelevant as the dataset lacks earlier data.
electricityBasePrice	Deprecated since 2020; irrelevant.
energyEfficiencyClass	Deprecated since 2020; irrelevant.
date	Redundant, as it duplicates date_full .
description	Redundant, as it only repeats information already provided in other columns.
facilities	Redundant, as it only repeats information already provided in other columns.
totalRent	Artificial column, a sum of baseRent , heatingCosts , and serviceCharge ; unnecessary for analysis.
newlyConst	Artificial column, constructionYear >= 2017 ? TRUE : FALSE

Upon removing the irrelevant fields from the Berlin dataset, the remaining fields include:

Data Analysis

To understand the contents of the columns in our dataset and how they interact, we visualize them. We will focus on how the districts are distributed across our dataset and how attributes, such as the **baseRent**, are distributed within them. Additionally, we will examine the apartments in our dataset and investigate whether there are any correlations present.

District-Level Apartment Overview and Model Considerations

As seen in the data overview chapter, our dataset does not include the districts directly but only the postal codes (`geo_plz`). To derive the districts, we join our dataset with a reference dataset that maps Berlin’s postal codes to districts. It is important to note that some postal codes belong to multiple districts. For simplicity, we associate each postal code with only one district. This simplification introduces a bias towards districts which appear earlier in our data, The following figure displays the frequency of each district represented in our dataset.

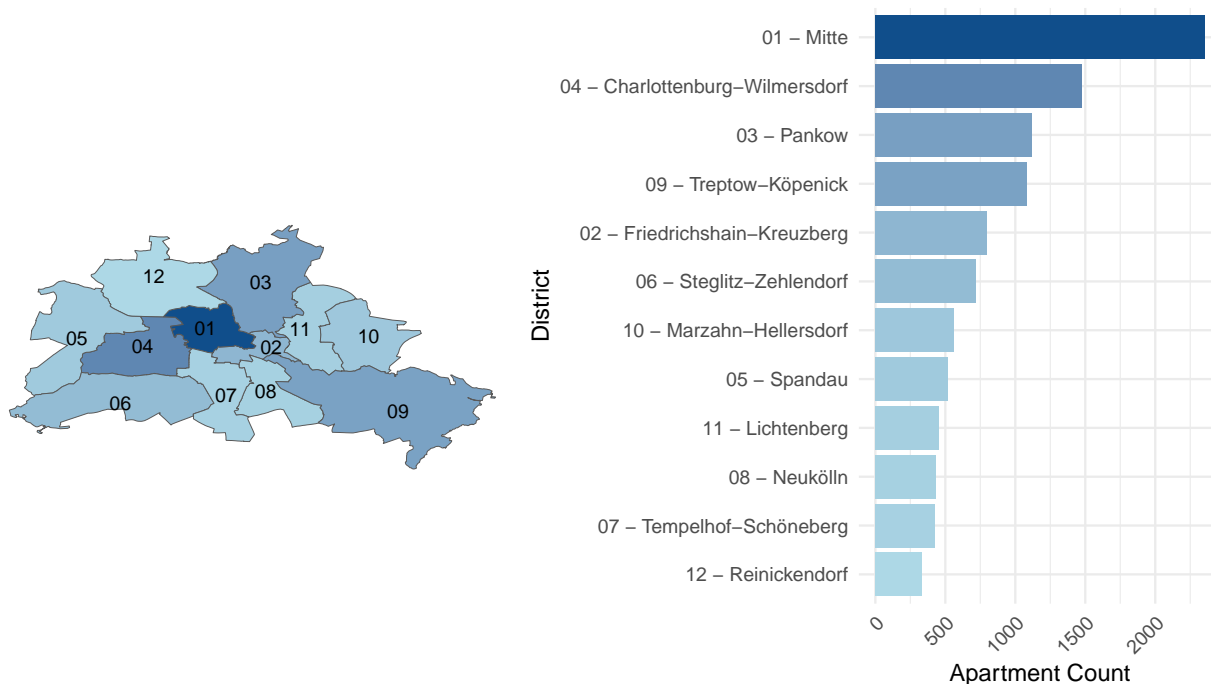


Figure 1: Apartment Count by District in Berlin

We observe that districts in central and eastern parts of Berlin are more frequently represented in the dataset compared to those in the southern or peripheral areas. Based on this we will evaluate later whether our model is bias towards districts that appear more frequently.

Furthermore, we analyzed the median prices across different districts. Our initial assumption was that districts closer to the city center have a higher median price compared to those in the outskirts. This assumption is based on the knowledge that central areas typically offer superior infrastructure and greater access to amenities and activity hubs.

Our initial assumption is confirmed: districts closer to the city center generally have higher rental prices, with Mitte standing out with several significant outliers. These outliers could pose challenges for our model if it lacks the robustness to accommodate them. Interestingly, when comparing this plot with the earlier analysis of apartment counts per district, the correlation between rental prices and the number of apartments appears weak. For example, Treptow-Köpenick, despite having a high number of apartments, shows relatively low rental prices.

Analysis of Rental Characteristics

The next step in understanding our dataset is to analyze characteristics of the apartments. This gives us an initial insight into the typical apartment types in our data and helps identify anomalies. First, we examine the distribution of the number of rooms across the apartments.

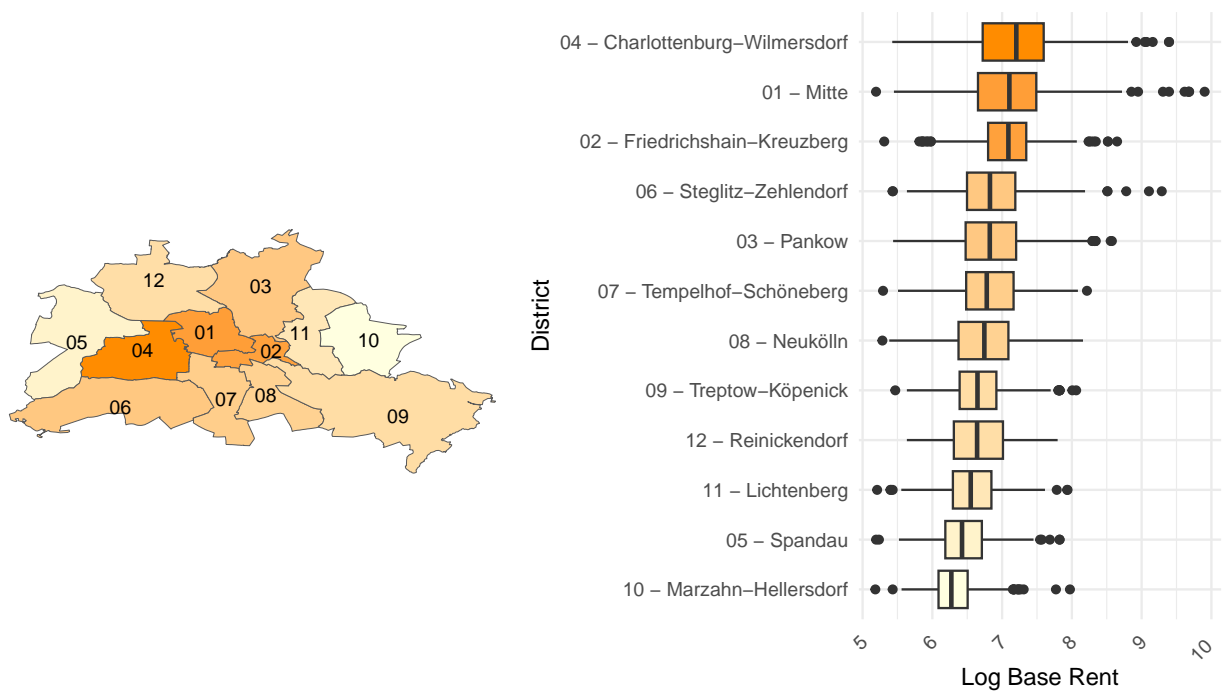


Figure 2: Median Price by District in Berlin




Figure 3: Count of Apartments by Number of Rooms

From the distribution, it is evident that the vast majority of apartments have a reasonable number of rooms, ranging between 1 and 8. However, there is one significant outlier, an apartment with an astonishing 100 rooms. While most of the data falls within expected ranges, such an outlier could have a disproportionate influence on our analysis and modeling.

One of our hypotheses about this dataset is that there should be a strong positive correlation between the number of rooms and the living space. Intuitively, as the number of rooms increases, the total living space is also expected to increase. Similarly, for smaller apartments, fewer rooms should correspond to smaller living spaces.

To visualize the relationship between living space and the number of rooms, we decided to exclude room counts that appear fewer than 10 times in the following visualization. This threshold ensures that the data we analyze is statistically meaningful. As a result, we focus on apartments with room counts greater than 8 for the following boxplot.




Figure 4: Living Space Distribution by Number of Rooms

The filtered boxplot confirms this hypothesis. Even for a room count like 8, which only has 16 occurrences, the trend remains consistent. Larger apartments have larger living spaces, which aligns with our expectations. A noteworthy observation is the prevalence of outliers in apartments with up to 5 rooms, especially for those with 3 and 4 rooms. These outliers are highly dispersed, potentially indicating that some of these entries might represent office spaces rather than standard apartments.

Another important aspect of our dataset is the condition of the apartments. This feature provides qualitative information about the state of the properties, with categories like *mint_condition*, *refurbished*, and *renovated*. However, the definitions of these categories are not immediately clear, particularly the distinction between “mint_condition” and the other classifications.

To better understand these categories, we compare the condition of the apartments to their rent values. This comparison allows us to infer potential meanings or differences between the conditions.

A clear order emerges among the apartment conditions. Apartments classified as *mint_condition* appear to represent the best quality. Similarly, *refurbished* properties seem to have a better condition than *renovated*.




Figure 5: Base Rent by Conditions

This ordering aligns with the intuition that better maintained or new apartments tend to have higher rents. Outliers might represent unique cases, such as misclassifications, unusual rental agreements or exceptional property features.

Another observation is that the missing entries (NA) have a wide interquartile range and contain several outliers, which indicates that these entries are spread across categories and may not represent a single uniform condition. This variability could introduce noise into our analysis and modeling. To address this, we aim to impute the missing condition values.

Exploring Correlations Between Variables

To better understand the relationships between variables in our dataset, we analyzed the correlations between them, focusing on how the predictors relate to the **baseRent**. This helps us identify expected patterns and uncover any unexpected relationships.

From the correlation analysis, we observed several expected relationships:

- Apartments with a high **baseRent** tend to have:
 - Higher **serviceCharges**, **heatingCosts** (additional rental costs)
 - Larger **livingSpace**, more rooms (**noRooms**)

These findings align with our expectations, as more expensive apartments generally offer more space and additional services.

One initially confusing correlation was between **thermalChar** and **yearConstructed**. After further research, we found that **thermalChar** represents the apartments energy efficiency level, which is often linked to the buildings construction year. Newer buildings tend to be more energy-efficient due to modern construction




Figure 6: Correlation Matrix

standards and regulations.

Interestingly, this energy efficiency only shows a weak negative correlation with **heatingCosts**. Contrary to our initial expectations, more energy-efficient apartments appear to incur slightly higher heating costs. One possible explanation is the cost disparity between energy sources, in 2020 and even today, dirtier energy sources are often more affordable than renewable or sustainable alternatives.

Imputation

Missing values in the data can significantly impact model performance, especially for Support Vector Machines (SVMs). While tree models such as those in **rpart** can handle missing values, SVMs from the **e1071** package cannot. SVMs simply exclude rows with missing values, which can lead to a scenario where the model is trained on as few as 10 observations.

In our initial data overview we identified the number of missing values in each column. To improve our understanding, we will visualize the top 10 columns with the highest percentage of missing values. These columns are also potential candidates for exclusion if we decide to reduce our model size. The following plot displays the absolute number of missing values out of a total of 10,304 entries.

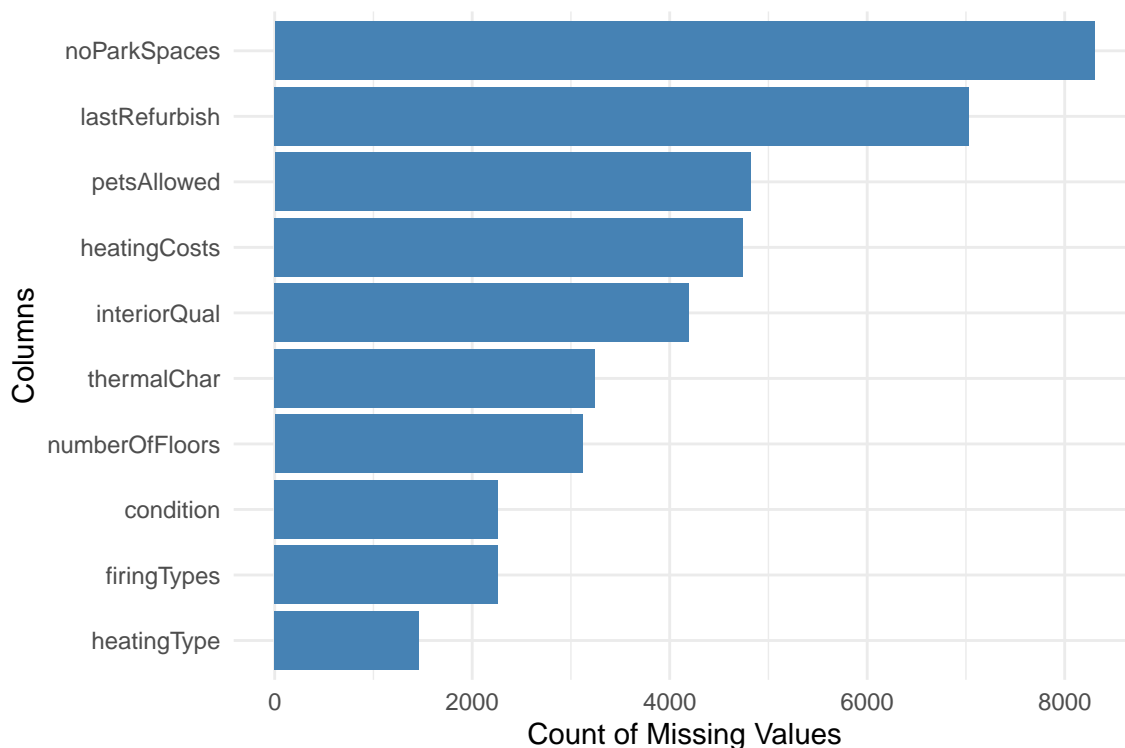


Figure 7: Top 10 Columns with Most Missing Values

For some columns, such as **lastRefurbished**, missing values have specific meanings. For example, indicating that the apartment has never been renovated. For the remaining columns we assume that the missing values represent missing data.

To address the values in **lastRefurbished** we will impute missing values with the year of construction, as this represents the most plausible time for the first or last renovation.

For all other columns we will use *Gibbs sampling* to impute values. This approach preserves the underlying distribution of variables and minimizes information loss during model fitting.

The imputation process using the *MICE* algorithm will generate five different datasets. For model training

we will build models on all five datasets and aggregate their results to produce a final output. For simplicity in visualizations, such as decision tree plots, we will use the first dataset generated by Gibbs sampling.

Model Overview

The following chapter begins with an overview of the machine learning algorithms we plan to use for our models. The content in this section is primarily based on our lecture slides, with citations provided for any external sources. Subsequently, we will train various models and evaluate them.

Regression Trees

Tree models are supervised learning algorithms used for both regression and classification. They work by splitting the dataset into smaller subsets based on decision rules derived from the features and forming a tree structure. Each split improves the prediction accuracy of the model on the training data. The process continues until a stopping criterion, such as minimum number of observations in a node or maximum tree depth, is met. The tree's leaf nodes represent the final predictions. For regression trees, this prediction is typically the average of the target variable within that leaf.

Regression trees are particularly useful for predicting continuous values as they can easily capture non-linear relationships. In our case, we aim to predict **baseRent** using a regression tree model.

The following procedure shows how regression trees work:

1. Initialization:

The process begins with the entire dataset. The mean of the target variable Y serves as an initial prediction, as it minimizes the **Residual sum of squares (RSS)**:

$$RSS = \sum_{i=1}^n (y_i - \hat{y})^2,$$

where \hat{y} is the mean of Y .

2. Splitting the data:

At each node, the dataset is split into two regions R_1 and R_2 , based on feature X_j and threshold x :

$$R_1 : X_j < x, \quad R_2 : X_j \geq x.$$

The goal is to find the split that minimizes the RSS for the two regions:

$$RSS = \sum_{i \in R_1} (y_i - \hat{y}_1)^2 + \sum_{i \in R_2} (y_i - \hat{y}_2)^2,$$

where \hat{y}_1 and \hat{y}_2 are the mean values of Y in R_1 and R_2 , respectively.

3. Recursive splitting:

This process of splitting is repeated recursively. At every step, the algorithm evaluates all possible splits for every feature and selects the one that minimizes the RSS.

4. Stopping criteria:

If we continue splitting until each node contained only a single observation, then each node would perfectly predict the value of its observation, leading to an overfitted model with an RSS of 0.

That is why we stop the tree growth when a predefined stopping criterion is met. This could be:

- The minimum number of observations in a node is reached.

- The maximum depth is reached.
- Further splits fail to reduce the RSS significantly.

5. Pruning with penalized least squares:

Overfitting is a common problem with regression trees when they grow too large. To address this, the tree is pruned using **Penalized least squares (PLS)**:

$$PLS(\alpha) = \sum_{j=1}^{|T_\alpha|} \sum_{i \in R_j} (y_i - \hat{y}_j)^2 + \alpha |T_\alpha|,$$

where:

- $|T_\alpha|$ is the number of terminal nodes (leaf nodes) in the pruned tree T_α ,
- $\alpha \geq 0$ is a complexity parameter controlling the trade-off between tree size and prediction accuracy.

The parameter α is chosen by using cross-validation, which ensures the pruned tree balances model complexity and performance.

6. Choosing the best subtree:

To find the optimal tree:

- The full tree is grown.
- The pruning process evaluates subtrees by varying α and calculates the relative error using cross-validation.
- The subtree corresponding to the smallest cross-validation error is selected.

Ensemble Methods

To understand how tree ensemble methods work, we need to know what bootstrapping is. Bootstrapping is a statistical technique that involves sampling with replacement from a dataset to create multiple subsets. Each bootstrap sample is the same size as the original dataset, but some observations may appear multiple times while others may be excluded. This technique is commonly used to generate diverse training subsets or to augment small datasets. In ensemble methods like bagging and random forests, bootstrapping enhances generalization and robustness by reducing variance only slightly increasing bias.

Bagging

Bagging is an ensemble learning technique that enhances accuracy and reduces variance by combining the predictions of multiple models. It involves training several instances of the same machine learning algorithm on different subsets of the data and then aggregating their predictions to form a final output. For regression tasks, this typically involves averaging the predictions. For classification tasks, the class with the majority vote is selected.

Bagging reduces model variance by averaging out individual errors, which leads to more robust and reliable predictions. This technique is especially useful when using base models that are prone to high variance, such as decision trees. An ensemble of decision trees is commonly referred to as a *forest*.

Random Forests

Random forests enhance the concept of *forests* by introducing additional randomness during the tree building process. At each node split, only a random subset of features is considered, which prevents the model from relying too heavily on the same dominant predictors. This also allows less important predictors to capture

more subtle patterns in the data. Typically, the number of features m considered at each split is given by $m = \lfloor \sqrt{p} \rfloor$, where p is the total number of features. This approach helps to reduce overfitting caused by a few strong predictors, which leads to more robust and reliable predictions.

Boosting

Boosting is an ensemble learning technique in machine learning that aims to improve the performance of weak models, such as a small tree (called *stump*) by combining their predictions to create a strong predictive model. The core idea is to iteratively train a series of models, where each new model focuses on correcting the mistakes made by the previous ones.

The process begins with a simple model, like a stump, which makes predictions on the training data. The instances that were misclassified or had large errors are then given higher weights, so that subsequent models pay more attention to correcting these errors.

In the next iteration, a new model is trained on the weighted data, where the misclassified have higher weights. After training, it is added to the ensemble and its predictions are combined with those of the previous models. The respect to errors the weights are updated again. This process repeats for several iterations, with each new model gradually improving the overall predictions.

The final prediction is made by aggregating the predictions of all models. In algorithms like AdaBoost, this is done by weighted voting, where better performing models have more influence on the result.

Boosting helps to reduce bias and variance. Initially, a simple model may have high bias and low variance, but as boosting progresses, the bias decreases, which improves the accuracy of our predictions. However, the increased flexibility and strong focus on errors can also lead to higher variance, making the model more sensitive to noise. To prevent overfitting regularization techniques and early stopping are often used.

We want to test the following boosting algorithms on our dataset:

- **Gradient Boosting:** Focuses on minimizing a loss function by fitting new models to the residuals of previous predictions.
- **XGBoost:** An optimized version of gradient boosting that enhances speed, performance and uses regularization techniques.

Support Vector Machines

Support vector machines (SVMs) are supervised learning algorithms that can be used for both classification and regression tasks. They work by finding a hyperplane that best separates data points into different classes. SVMs excel in situations where relationships between variables are complex and potentially non-linear. This chapter provides an overview of how SVMs work, with a focus on their extension to regression tasks and the kernel trick, which enables handling non-linear relationships effectively.

Linear Support Vector Machines

In classification SVMs aim to identify a hyperplane $f(x)$ that satisfies:

$$\hat{y}(x) = \text{sign}(f(x)), \quad f(x) = \beta_0 + \sum_{j=1}^p \beta_j x_j,$$

where the hyperplane separates the data into distinct classes while maximizing a margin, which is the distance between the hyperplane and the closest data points from each class, known as *support vectors*.

However, real world datasets often contain overlapping classes, making perfect separation impossible. To address this SVMs use a *soft margin* that allows some misclassifications. The trade off between margin width and classification errors is controlled by the penalty parameter C . The optimization problem for a

soft margin SVM is:

$$\text{maximize } M, \quad \text{subject to } y_i f(x_i) \geq M(1 - \xi_i), \quad \xi_i \geq 0, \quad \sum_{i=1}^n \xi_i \leq C,$$

where ξ_i are slack variables representing the extent of misclassification.

While linear SVMs are effective for linearly separable data, they struggle with non-linear relationships, where a simple hyperplane cannot adequately separate the classes.

Non-linear Support Vector Machines

To overcome the limitations of linear SVMs, data can be projected into higher-dimensional feature spaces, where a linear hyperplane can separate the data. This involves mapping the input x from its original space \mathbb{R}^p to a higher-dimensional feature space \mathbb{R}^q , where $q > p$:

$$\phi(x) : \mathbb{R}^p \rightarrow \mathbb{R}^q.$$

In this transformed space, a linear decision boundary may successfully classify the data. However, explicitly computing $\phi(x)$ is computationally inefficient, especially for large datasets or high-dimensional mappings.

The *kernel trick* addresses this challenge by avoiding explicit computation of $\phi(x)$. Instead, SVMs calculate the inner product in the feature space $\langle \phi(x_i), \phi(x_j) \rangle$, indirectly using a kernel function $K(x_i, x_j)$. This enables efficient computation in the higher-dimensional space without ever explicitly performing the transformation. The decision function then takes the form:

$$f(x) = \beta_0 + \sum_{i \in S} \alpha_i K(x, x_i),$$

where S represents the set of support vectors that define the hyperplane, enhancing computational efficiency by focusing only on these critical points. The values α_i are the Lagrange multipliers corresponding to each support vector x_i . These multipliers are determined during the training process and reflect the importance of each support vector in defining the optimal hyperplane.

By leveraging the kernel trick, SVMs maintain computational efficiency while gaining the flexibility to tackle complex, non-linear problems.

Regression with Support Vector Machines (Rasifaghihi, 2023)

Support vector regression (SVR) is an extension of SVMs for regression tasks. The goal of SVR is to identify a function $f(x)$ that accurately predicts the target variable, while minimizing model complexity to enhance generalization and reduce overfitting.

A distinctive feature of SVR is the introduction of an ϵ -insensitive tube around the regression hyperplane. This tube defines a tolerance margin within which deviations between predicted and actual target values are not penalized. This creates a balance between model complexity and its ability to generalize effectively.

For a linear function $f(x) = wx + b$, this means to minimizing w , which makes it less sensitive to small variations in the input data.

The SVR problem is formulated as a convex optimization task with the following objectives:

1. Ensure that the deviations between predictions $f(x)$ and actual target values y remain within ϵ for all training data points.
2. Minimize the complexity of the function $f(x)$ by keeping w small.

In cases where the optimization problem is infeasible due to noisy data or overlapping regions, *slack variables* (ξ_i^+ and ξ_i^-) are introduced. These variables allow for deviations beyond the ϵ -insensitive tube, representing data points that fall outside the tube.

The optimization problem can be expressed as:

$$\min \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n (\xi_i^+ + \xi_i^-),$$

subject to the following constraints:

$$y_i - f(x_i) \leq \epsilon + \xi_i^+, \quad f(x_i) - y_i \leq \epsilon + \xi_i^-, \quad \xi_i^+, \xi_i^- \geq 0.$$

The regularization strength is inversely related to C , meaning higher values allow the model to prioritize accuracy at the expense of simplicity.

SVR assigns zero prediction error to points within the ϵ -insensitive tube. For points outside the tube (slack variables), the penalty is proportional to the magnitude of their deviation. The tolerance for small deviations and penalties for larger errors make SVR more robust to overfitting.

Model Implementation

This chapter describes the implementation of the models trained on our dataset, outlining key considerations and decisions made to optimize our models. The following topics are covered: model training, hyperparameter optimization and managing outliers.

One challenge we faced was the diversity of geographic data inside the `geo_plz` variable, which includes 214 unique levels across 10,000 rows. This introduces the potential issue of encountering postal codes in the test data that were not present in the training data. To address this we removed the `geo_plz` variable by grouping data into broader districts and subdistricts within the `regio3` category. This also helps to capture meaningful geographic relationships while minimizing overfitting risks.

The implementation of all basic models is included in `eric_model.R`. Hyperparameter optimization (HPO) is performed exclusively on a single Gibbs-imputed dataset using 10-fold cross-validation in `eric_hpo.R`. This approach was selected to save computational resources, as Gibbs iterations maintain consistent distributions and variables with the highest correlations in our dataset have fewer missing values.

All models are then trained using the hyperparameters determined through HPO with an 80/20 train-test split on all of our imputed datasets to evaluate their performance. Additionally, outlier management is an aspect of our analysis, with experiments conducted both including and excluding outliers. Outlier free datasets are reserved for a final comparison. Finally, we will explore if it is possible to reduce the model size and evaluate how the models perform with fewer columns and potentially less noise. The following table summarizes the results of this chapter:

Table 2: Comparison of Model Performance

Method	Average MSE	Average RMSE	Average MAE
Regression Tree	129592	360	208
SVM	114781	339	170
Random Forest	92827	305	167
Gradient Boosting	92215	303	170
XGBoost	73161	270	154

Regression Trees

In this section, we focus on training a **regression tree** using the `rpart` package and visualizing the results. To determine the optimal depth or number of splits for our tree, we use cross-validation. The `rpart` package simplifies this process by calculating the cross-validated error for each of the subtrees. These cross-validation

results are stored in a complexity parameter table (`cptable`), which provides the cross-validation errors for the subtrees. This table helps us to assess the complexity needed when we prune our tree. Below, we display the first 10 rows of the `cptable` for our tree:

Table 3: Complexity Table for Tree

CP	nsplit	rel error	xerror	xstd
0.40	0	1.00	1.00	0.09
0.13	1	0.60	0.60	0.07
0.08	2	0.47	0.52	0.05
0.05	3	0.39	0.43	0.05
0.03	4	0.34	0.38	0.03
0.02	5	0.31	0.37	0.03
0.02	6	0.29	0.35	0.03
0.01	7	0.27	0.33	0.03
0.01	8	0.26	0.32	0.03
0.01	9	0.25	0.32	0.03

From the `cptable`, we can use two approaches to prune the tree:

1. **Lowest Cross-Validation Error:** Prune the tree to the complexity parameter corresponding to the lowest cross-validated error (`xerror`).
2. **Standard Deviation Rule:** Use the complexity parameter with the simplest tree whose error is within one standard deviation of the minimum cross-validated error. This approach is recommended by the authors of the `rpart` package.

The following plot illustrates the first complexity parameters and cross-validation errors for our tree:

Using these methods, we pruned the full tree. The first splits of the pruned tree are shown below:

The pruning process helps to create a more interpretable tree while maintaining or even enhancing its predictive performance. Below, we compare the performance of the full tree and the two pruning techniques on the test data, averaged across all five imputed datasets. Note that the *full tree* does not necessarily imply only one node per leaf, other stopping parameters influence the structure. Having one terminal node per observation would be excessive and unnecessary.

Table 4: Comparison of Tree Performance

Method	Average MSE	Average RMSE	Average MAE
Full Tree	130516	361	205
Best xError	129592	360	208
Best (Authors' Method)	165023	405	254

From the table we observe the following:

1. **Full Tree:** Surprisingly, the full tree performs well on the test dataset, achieving competitive results in MSE and RMSE and the best results for mean absolute error. This suggests a high degree of variance in the dataset.




Figure 8: Default Tree Complexity




Figure 9: Tree First Splits

2. **Best xError:** This pruning technique achieves significantly better results compared to the method recommended by the authors of the `rpart` package, as it has the lowest MSE and RMSE out of all methods.
3. **Authors Method:** While this approach prioritizes simplicity and generalization, it results in higher error values compared to the other methods. This means that this pruning method may not be the optimal choice for this dataset.

Support Vector Machines

In this section, we focus on training a **SVM** for regression using the `e1071` package. SVM Regression is a distance based algorithm, which means they calculate distances between data points to make decisions. One key consideration when using SVMs is that variables with larger scales can dominate the distance calculations, potentially leading to inaccurate results. This is because the SVMs decision boundary relies on calculating the distance between points. Therefore normalizing or standardizing data is crucial to ensure that each feature contributes equally to the model.

For tuning the SVM parameters and validating its performance, we performed cross-validation on one of our imputed datasets with the following settings, where the best parameter combination is indicated in bold.

- **Kernel Function:** (`polynomial`, `radial`, `sigmoid`, `linear`)
The kernel function defines the type of decision boundary the SVM will create. It maps the input features to a higher-dimensional space where linear separation is possible.
- **Cost (C):** (0.5, **1**, 2, 4)
The cost parameter controls the trade-off between minimizing training error and ensuring good generalization to new data. It penalizes misclassification of data points. A higher cost leads to a more complex model that fits the training data closely, which may result in overfitting.
- **Gamma ():** (0.01, **0.03**, 0.5, 1, 2)
Gamma influences individual training points when using non-linear kernels (e.g., *radial*, *polynomial*). A high gamma value causes a potentially overfitted model. A low gamma value results in a smoother decision boundary that is less sensitive to individual data points. (Günay, 2021)
- **Degree:** (2, **3**, 6, 10)
The degree parameter applies only to the *polynomial* kernel. It specifies the degree of the polynomial used to compute the kernel. A higher degree increases the complexity which could lead to overfitting.
- **Coef0:** (0, 1, **5**, 10)
The coef0 parameter is relevant only for *polynomial* and *sigmoid* kernels. It controls the impact of the higher-dimensional feature space on the decision boundary. (Lejlot, 2013)

In this case the results of SVM fitting over all imputed datasets are as follows:

Table 5: SVM Performance

Method	Average MSE	Average RMSE	Average MAE
SVM	114781	339	170

When comparing SVMs to simple regression trees, we can see that SVMs performed slightly better. Decision trees are more interpretable, but also prone to overfitting, especially when there are many features or complex relationships in the data. SVMs tend to generalize better, which leads to lower error metrics. While both models have their advantages, the SVMs performance suggests it is better suited for this particular dataset, likely due to its improved ability to handle outliers and avoid overfitting. To evaluate these

statements, we plotted a regression line with the training data predictions for `livingSpace`, one of the variables with the highest correlation to `baseRent`. The plot indicates that the model demonstrates resilience to certain outliers and appears to generalize reasonably well, as its prediction is between the range of the minimum and maximum `baseRent`.




Figure 10: Living Space vs Base Rent

Random Forest

In this section, we focus on training a **random forest** using the `randomForest` package. Random forests in R cannot natively handle variables with more than 53 categorical features. This restriction posed a challenge in our dataset, as the `regio3` variable contains more than 53 unique categories. Consequently, we excluded this variable from the random forest training process. While this exclusion makes a direct comparison with other models somewhat unfair, we believe that random forests' ability to reduce noise and capture complex interactions still allows them to outperform single decision trees and SVMs in general.

To optimize the random forest model, we performed cross-validation on one of the imputed datasets using the following parameters, where the best parameter combination is indicated in bold.

- **Number of Trees:** (100, **500**, 1000)
The number of trees in the forest. A higher number of trees usually improves performance but increases computation time.
- **Maximum Features:** (2, 5, 10, **20**)
The number of features randomly selected at each split. A smaller value can help reduce overfitting (a common rule of thumb is to use \sqrt{n}), while a larger value allows the model to capture more complex relationships.

- **Node Size:** (2, 5, 10, 20)

The minimum number of observations in a terminal node. Larger values result in simpler trees and can reduces complexity.

The results of random forest fitting on all imputed datasets, excluding the `regio3` variable, are summarized below:

Table 6: Random Forest Performance

Method	Average MSE	Average RMSE	Average MAE
Random Forest	92827	305	167

To visualize the variable importance of our first model using imputed data, we generated a plot that illustrates the impact of each variable on the model based on the `incNodePurity`. A higher `incNodePurity` value indicates a greater influence of the variable on the random forest model.




Figure 11: Wordcloud of Titles

We can observe that variables with high correlations exert the most influence on the random forest. Based on these findings, we plan to reduce the models complexity in a later chapter and explore whether a smaller model, trained on less noisy data, could outperform models trained on the full dataset.

Random forests demonstrated superior performance in terms of error metrics compared to both SVMs and single decision trees. The ensemble approach reduced the overall noise in the data and captured complex interactions among variables more effectively. Although the exclusion of `regio3` may have slightly impacted its performance.

Gradient Boosting Machine

In this section, we focus on training a **gradient boosting machine (GBM)** using the `gbm` package in R. Unlike random forests, gradient boosting emphasizes optimization by minimizing a loss function iteratively, which often results in highly predictive models, but with a higher risk of overfitting.

To optimize the GBM model, we performed a grid search using cross-validation on one of the imputed datasets. The parameters explored and their implications are as follows:

- **Number of Trees** (100, 200, 300, **500**)
The number of boosting iterations. More trees typically improve performance up to a point but increase computational cost and the risk of overfitting.
- **Tree Depth** (1, 3, **6**, 9)
The maximum depth of each tree. A smaller depth helps prevent overfitting and ensures the model captures simpler interactions, while deeper trees can model more complex relationships.
- **Shrinkage** (**0.05**, 0.1)
Also known as the learning rate, this parameter controls how much each tree contributes to the final model. Smaller values slow down the learning process and can lead to more robust models.
- **Minimum Observations per Node** (**10**, 20)
The minimum number of observations in terminal nodes. Higher values result in simpler trees and prevent overfitting by reducing model complexity.

Using the best combination (**bold**), we evaluated the GBMs performance across all imputed datasets, summarized as follows:

Table 7: GBM Performance

Method	Average MSE	Average RMSE	Average MAE
GBM	92215	303	170

A partial dependence plot illustrates the relationship between a predictor variable and the target variable while holding other features constant. This visualization helps interpret the influence of individual predictors on the models predictions and provides insights into how a specific variable affects the target. (Molnar, 2024) The plot below highlights the relationship between `livingSpace` and `baseRent`.

The partial dependence plot for `livingSpace` demonstrates a good fit, considering the high correlation with the target variable. The absence of abrupt changes in the plot highlights the robustness.

The variable importance analysis for our GBM models reveals that the `regio3` column is significantly more influential than the `district` column. Since the `district` column is an abstraction of `regio3`, removing `regio3` results in a substantial loss of information. Additionally, an interesting observation is that `noRooms` holds greater importance in the GBM models compared to `yearConstructed`, which contrasts with the findings from the random forest models, where `constructionYear` is more influential.

Gradient boosting only outperformed single decision trees in all error metrics. Compared to random forests and SVMs, GBMs demonstrated slightly higher errors but offered the advantage of interpretable partial dependence plots. While random forests are more robust to noise due to their averaging mechanism, gradient boosting prioritizes correcting previous errors, which can lead to overfitting to noise. This tendency may explain why the model underperforms compared to random forests and SVMs when dealing with our data, which contains a high levels of noise.




Figure 12: Partial Dependence Plot




Figure 13: GBM Variable Influence

XGBoost

In this section, we focus on training an **XGBoost** model using the `xgboost` package in R. (Developers, 2022) XGBoost is an optimized implementation of gradient boosting that often leads to better and faster results.

Before training the model, it is essential to preprocess the categorical variables. XGBoost in R requires the input data to be in a numeric matrix form, so we need to *one-hot encode* or *label encode* the categorical variables. In our dataset none of the categorical variables have an inherent order, meaning they are nominal. Because of that we used **one-hot encoding** instead than *label encoding* to avoid introducing any unintended ordinal relationships.

To optimize the XGBoost model, we performed a *random grid search* over a range of hyperparameters. Given the computational cost of a full grid search, we implemented a random sampling strategy. The parameters explored are as follows:

- **eta** (0.01, 0.05, **0.1**, 0.3, 0.5)
The learning rate, which controls the contribution of each tree. A smaller eta can lead to better generalization but requires more trees to achieve convergence.
- **max_depth** (1, 3, **6**, 9)
The maximum depth of each tree. This parameter controls the complexity of the individual trees, with deeper trees modeling more complex relationships.
- **min_child_weight** (1, **2**, 5, 10)
The minimum sum of instance weight (hessian) in a child. It helps control overfitting by ensuring that trees grow only when a sufficient number of data points are present in a split.
- **subsample** (0.6, **0.8**, 1.0)
The proportion of rows used for training each tree. Lower values help reduce overfitting by introducing randomness, but too low may hurt performance.
- **colsample_bytree** (**0.6**, 0.8, 1.0)
The proportion of features used for training each tree. It helps prevent overfitting by limiting the model's reliance on any particular feature.

Using the best combination of hyperparameters, we evaluated the performance of XGBoost across all imputed datasets:

Table 8: XGBoost Performance

Method	Average MSE	Average RMSE	Average MAE
XGBoost	73161	270	154

We can observe that XGBoost outperforms GBMs and to better understand why, we examined the models variable importance. The variable importance analysis for the XGBoost model reveals slight differences compared to the GBM model. For example the **district** column holds more importance than the **regio3** column (subdistricts), suggesting that we can potentially remove **regio3** without significant loss of information.

XGBoost delivered the best results out of all methods tested, even with the default parameters. This demonstrates XGBoosts ability to capture the most variance in the data while avoiding overfitting and XGBoosts state-of-the-art model status.

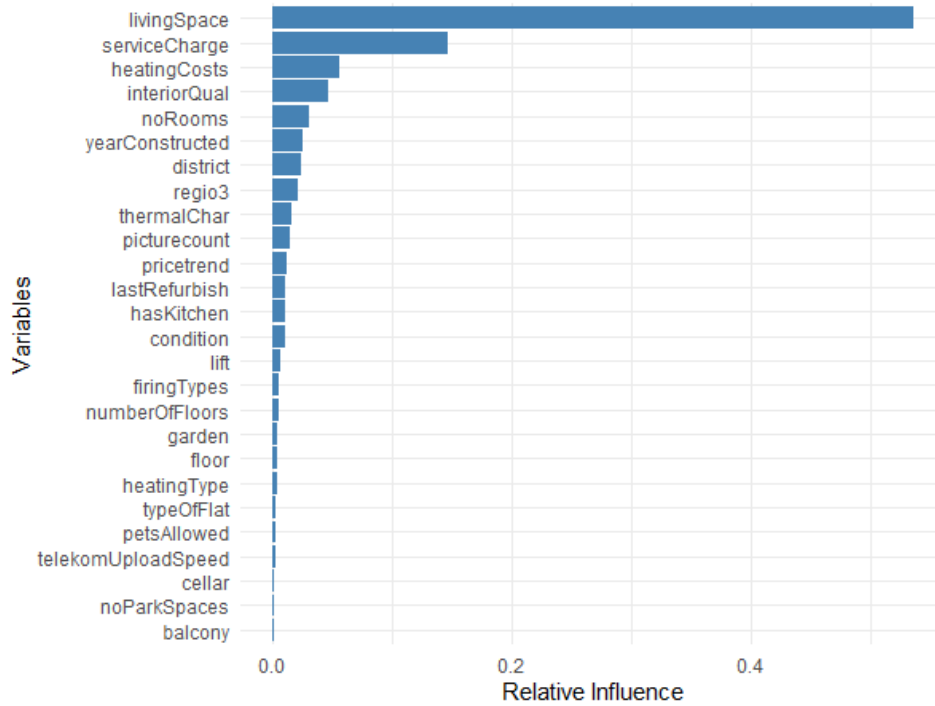


Figure 14: XGB Variable Influence

Outlier Handling

The dataset has a significant variance in prices, with few high-priced apartments that create a *long-tail* distribution. To assess the robustness of our models against outliers, we will begin by examining the true vs. predicted values for our best performing model, which was XGBoost.

The plot reveals that the model performs well for most non-outlier base rents, with predictions closely aligned to the true values. However, there is noticeable underprediction for base rents above 2,600, while higher base rents tend to have wider variation in predictions. Despite this, the model demonstrates some robustness against outliers.

To further evaluate the robustness of the model, we propose the following tests:

- 1. Training and Testing on Outlier-Free Data:**

Retrain the models on a dataset that excludes outliers and evaluate their performance on outlier-free test data.

- 2. Testing Existing Models on Outlier-Free Data:**

Use the previously trained models (which included outliers during training) to assess their performance on outlier-free test data.

If the models trained with outliers perform as well as those trained without outliers, it would indicate that our models are robust and not significantly influenced by outliers. Vice versa if the performance of the outlier-free models is superior, it suggests that excluding outliers could improve predictions for common apartments.

One alternative approach we will not consider in this report is training the model on the log-transformed `baseRent`. While this method effectively scales down the impact of outliers, it introduces additional complexity that may not align with the objectives.

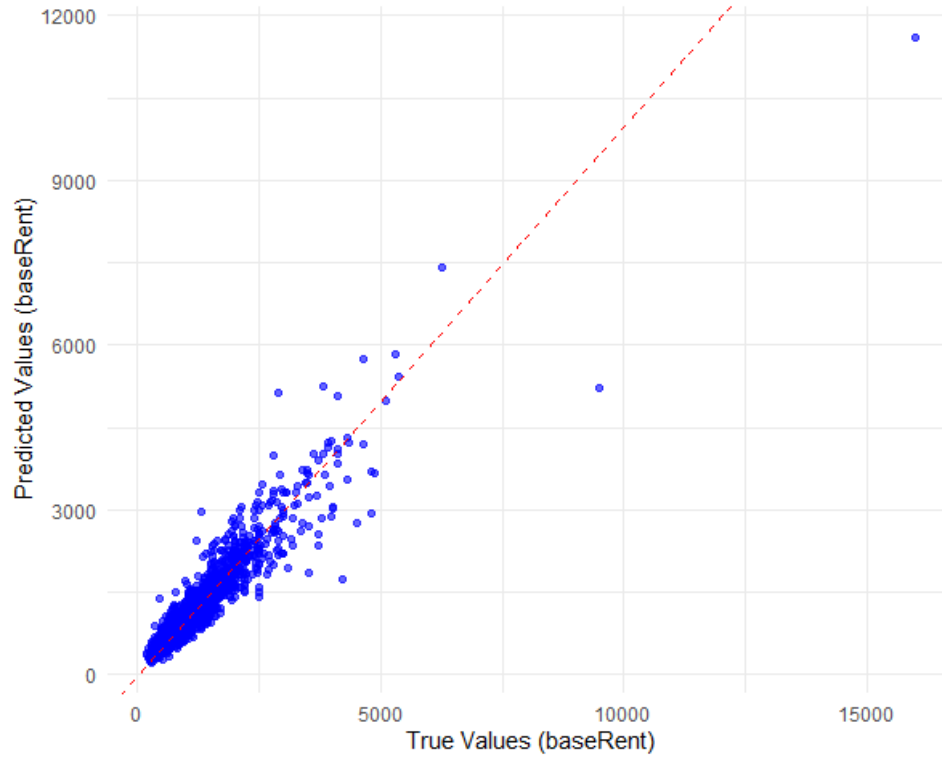


Figure 15: True vs Predicted Base Rent

There are several reasons we think it might be beneficial to get better predictions for more common apartments:

- Outliers may represent corrupted data that do not occur in real-world scenarios.
- Most potential users of the model may not be wealthy and are likely more interested in predictions for mid-range apartments.
- Prioritizing accuracy for mid-range apartments aligns better with the needs of the target audience, ensuring more practical and relevant predictions.

By focusing on this approach, we aim to enhance the reliability and usability of our predictions.

Outlier Removal

To detect outlier apartments, we use a boxplot logic. After this definition, all values bigger or smaller than 1.5 times the interquartile range from the median get excluded. The following plot shows the `baseRent` outliers for our whole dataset:

We can see from the boxplot, that there are only expensive apartments that are considered outliers. Therefore it is only necessary to calculate the upper bound to define it as an outlier border.

When calculating the upper bound, all flats with a price higher than 2646.25€ should be excluded.




Figure 16: Boxplot for Base Rent

Model Performance

Our initial hypothesis is that models that are better equipped to handle noise, such as random forests, will show smaller performance increases when trained with and without outliers, when the evaluation is done on outlier free test data. Models like Gradient Boosting, Regression Trees and XGBoost often benefit from the removal of outliers since their algorithms are sensitive to extreme values that can influence the structure of the decision boundaries or the boosting process. These models can more effectively generalize when the noise introduced by outliers is reduced.

It is important to note that we did not optimize the hyperparameters for the models trained on the dataset without outliers. All models were evaluated on outlier free data. The error metrics are presented in the table below:

Table 9: Comparison of Model Performance Metrics with/without Outliers

Method	Trained With Outliers	Average MSE	Average RMSE	Average MAE
Regression Tree	YES	63658	252	179
Regression Tree	NO	52322	229	168
SVM	YES	43701	209	143
SVM	NO	38546	196	138
Random Forest	YES	44653	211	146
Random Forest	NO	39227	198	139
Gradient Boosting	YES	46736	216	154
Gradient Boosting	NO	34438	186	135
XGBoost	YES	39065	198	134

Method	Trained With Outliers	Average MSE	Average RMSE	Average MAE
XGBoost	NO	32187	179	126

Here we can see that all methods significantly improved on the outlier-free data when we trained them without outliers. The order of the models, based on their improvement in RMSE (italic) after removing outliers, is as follows:

1. Gradient Boosting *30*
2. Regression Tree *23*
3. XGBoost *19*
4. SVM *13*
5. Random Forest *13*

This order could be due to the specific nature of each algorithms sensitivity to outliers and our selected parameters. For example, Gradient Boosting iteratively builds weak learners and is particularly impacted by outliers because each subsequent tree tries to correct for errors from previous ones. If the initial trees are influenced by outliers, it can lead to a cascading effect of poor learning.

Regression Trees, while also sensitive to outliers, tend to generalize better than weak learners in the presence of outliers, but extreme outliers can still influence splits.

XGBoost employs regularization techniques that can reduce the impact of outliers, but it still performs best when noise is minimized.

SVM and Random Forest, on the other hand, are less sensitive to individual data points but can still be affected by outliers. Random Forest, being an ensemble method, is somewhat more robust due to averaging over many trees, but it too benefits from cleaner data for optimal performance.

In summary, the degree of improvement in RMSE across different models when outliers are removed is likely related to the inherent characteristics of each model in handling outliers and noise. Models that rely more heavily on learning from the Variance, such as Gradient Boosting, show the greatest improvement when outliers are removed.

References

- Bar, C. (2020). Apartment rental offers in germany. In *Kaggle*. <https://www.kaggle.com/datasets/corrieaar/apartment-rental-offers-in-germany/data>
- Developers, X. (2022). XGBoost tutorial. In *XGBoost Documentation*. <https://xgboost.readthedocs.io/en/latest/R-package/xgboostPresentation.html>
- Günay, G. (2021). Understanding parameters of SVM. In *Kaggle*. <https://www.kaggle.com/code/gorkemgunay/understanding-parameters-of-svm>
- Lejlot. (2013). Scikit-learn SVC coef0 parameter range. In *Stack Overflow*. <https://stackoverflow.com/questions/21390570/scikit-learn-svc-coef0-parameter-range>
- Molnar, C. (2024). A guide for making black box models explainable. In *Interpretable Machine Learning*. <https://christophm.github.io/interpretable-ml-book/pdp.html>
- Rasifaghihi, N. (2023). From theory to practice: Implementing support vector regression for predictions in python. In *Medium*. <https://medium.com/@niousha.rf/support-vector-regressor-theory-and-coding-exercise-in-python-ca6a7dfda927>